



CLOUD COMPUTING SECURITY

MECR1043

SECTION 01

LECTURER

DR. MOHD ZAMRI OSMAN

FINAL PROJECT

EFFECTIVENESS OF ADVANCED THREAT MITIGATION STRATEGIES IN SECURING MULTI-CLOUD ENVIRONMENTS

<https://youtu.be/Fd1Md9EvHFQ>

NAME	MATRIC NO.
KEERTENNAH DEVI A/P PONNAMBALAM	MEC254026
SHWETHA A/P C.RAVICHANDRAN	MECS254010
NETRHA SHRI A/P RAMA MOORTHY	MEC254040

1.0 Introduction

The rise of the multi-cloud is a significant indication of a change in the way in which companies are now creating their infrastructure in the cloud. Deploying the workload across various cloud infrastructures instead of completely investing in one service provider is the most commonly used model to mitigate vendor lock-in risks, the reliability of the services that are provided, as well as the sustainability of the organization. While there are obvious economic benefits to these factors, they create a far more complex scenario compared to a single-cloud setup when it comes to security.

Heterogeneity can be highlighted as an essential aspect of multi-cloud infrastructure. The dissimilar designs of identity management, access control models, cloud provider-level authorization policy definitions, monitoring infrastructure, and logging infrastructure provided in multi-cloud environments enhance avenues to facilitate cloud security. Consequently, security professionals must contend with a portfolio of cloud security control plane instances concurrently and in most circumstances without any level of interoperability between these different instances. Heterogeneity involves not only costs in multi-cloud environments but also weaknesses that can potentially serve as a target vulnerability to security threats where different infrastructure fails to integrate in a cohesive environment.

Furthermore, statistics from the industry have previously shown that the risks of being compromised, as well as the slow compromises in the system, pose a greater threat relative to the complex infrastructures used in breaches in the system. This has resulted in the complexity of the threats being more pronounced, especially considering the idea of the mixed models used in the system, as well as the fact that some threats in the system could also pose a threat in the usage of other similar systems in the environment. A critical assessment of the underlying threat vectors has thus been required before the assessment of the threats from the prevailing advanced threat mitigation measures in the system.

Among the most significant security challenges within the threat environment of the multi-cloud, the focus of this section will be on three aspects of the security environment, namely, identification-based attacks, policy inconsistencies with regard to multi-cloud environment, and the issue of the absence of visibility in the multi-cloud environment.

2.0 Identity-Based Attacks in Multi-Cloud Environments

2.1 Identity as the Dominant Security Control Plane

Security enforcement has transformed profoundly as a direct consequence of cloud computing. In this context, there was an inherent move from network-based perimeter environments to identity-based environments. Practically every interaction in cloud environments involves identity and access management systems, which include processes such as automation, service and user authentication. Considering identification as the essential basis for trust-based determinations, as opposed to network location or infrastructure access, current concepts or ideas of security such as Zero Trust formalize this shift (NIST, 2020; Rose et al., 2020).

In multi-cloud systems, several IAM models operate simultaneously with variations in terms of role structures, authorization mechanisms, and authentication methods. It is, therefore, becoming increasingly challenging. In order to reduce management burdens and enhance usability, organizations may turn to federated identity models, where one identity provider uses its services to authenticate numerous other cloud services. This on the other hand reduces management burdens but focuses risk in a centralized area.

Identity is a high-value target from a security perspective. Leaked credentials allow attackers to operate within the bounds of authorized access. This risk is heightened in multi-cloud systems, as federated trust relationships can allow attackers to use one compromised identity to authenticate across different platforms. This consolidation of trust makes containment more challenging and vastly increases the blast radius of identity compromise.

Pivotal in this argument is the fact that identity centralization inherently increases security, which altogether dismisses the level of maturity required at the level of governance to manage federated trust correctly. Without continuous verification, strict role scoping, and consistent cross-cloud management, identity will become a systemic point of failure, not a security control, which enables compromise to spread between cloud platforms instead of remaining isolated.

2.2 Credential Abuse and Cross-Cloud Lateral Movement

One of the first access vector approaches which is most commonly observed in cloud security incidents remains credential misuse. To procure authentic authentication material, attackers often resort to phishing attacks, exposed API keys, misuse of authentication tokens like OAuth, and unsafe credential storage procedures (Mandiant, 2023; Verizon, 2024). By avoiding traditional perimeter security controls as well as security detection systems, an attacker may establish a direct API communication once the credentials have been compromised.

Credential abuse enables the possibility of cross-cloud lateral movement through federated identity interactions in multi-cloud environments. This could be due to overly permissive trust configurations or poorly monitored federation protocols, whereby attackers can move from one compromised environment to another without extra exploitation steps. This, in turn, presents a possibility that would breach one platform which may spread throughout the entire cloud ecosystem.

Besides, inconsistent logging and identification alerts across cloud providers further complicates the detection of such activities. The differences in terms of telemetry granularity, logging formats, alert levels, and authentication anomalies may be flagged in one environment but ignored in another. Attackers utilize this asymmetry by intentionally shifting activity onto less-monitored platforms to increase dwell time and decrease the likelihood of detection.

Credential exploitation in multi-cloud reveals the mistaken belief that cloud platforms function as separate security domains from a fundamental standpoint. Since it blurs these lines, credential breach poses a systemic risk; therefore, comprehensive cross-cloud threat modelling is required.

2.3 Privilege Sprawl and Non-Human Identities

Another issue associated with cloud computing is privilege sprawl, where the issue is magnified when implementing the multi-cloud environment due to the tendency of cloud administrators to over-grant access to ensure conformance while minimizing administrative hurdles associated with semantic role variances as well as the permission variances across

different cloud environments (Behl & Behl, 2017). Accumulation of identity leads to privilege model violations of the principle of least privilege.

However, the biggest challenges occur when we have non-human identities like automation credentials, application roles, or service accounts. The construction of the DevOps pipelines, the automation of the infrastructures, as well as the services, depend on these non-human identities, sometimes being long-lasting, not properly documented, and seldom assessed. The biggest challenges occur due to the diverse tool sets, and the visibility problems while managing these non-human identities in the environment of the hybrid cloud (CISA, 2023).

Furthermore, an attacker may focus on non-human identities because these may not necessarily have proper authentication restrictions such as multi-factor authentication and do not receive the same level of concern as human identities. Persistent and undetected access to multiple cloud platforms can be gained if these identities are compromised, mostly if the credentials have not changed or were not well monitored.

Most importantly, privilege sprawl is a reflection of a compromise in organization that prioritizes the pace of operations over the robustness of security. Such a balance introduces risk in a multi-cloud environment as the automation identity extends as a lasting attack vector and the concept of over-privilege is elevated to routine.

3.0 Policy Inconsistency Across Cloud Platforms

3.1 Fragmentation of Security Policy Models

Security policies define how network segmentation, encryption, logging, and access control are done in the cloud environments. These policies must be implemented consistently for different providers using different syntax, abstractions, and logics of enforcement in multi-cloud (ENISA, 2021). Due to the differences in the default behaviour and assessment methodologies, inconsistent security efficacy results in policies sharing the same intention.

Security teams often have to manually translate provider-specific parameters into organizational security requirements. This process introduces ambiguity and opens up the

opportunity for misinterpretations, especially in dynamic environments where policies are changing regularly. Besides, it is operationally complicated to maintain consistent compliance across various platforms, especially as settings become larger.

Instead of focusing on the most critical asset, fragmented policy models further provide an attacker with the opportunity to capitalize on the least effective setup. Taking into account that policy inconsistencies increase the probability of misconfiguration, attackers prefer to target environments with low policy enforcement and misaligned policy.

Significantly, the fragmenting policies represent a failure in governance rather than a technology constraint. When the entire security posture is hindered just because the solutions for individual platforms undermine the security purpose, the end solution puts the responsibility in the hands of the service providers rather than the risk appetite of the company.

3.2 Configuration Drift and Exploitable Weaknesses

The issue of configuration drift, where security configurations continuously move away from the desired baseline, is due to policy inconsistencies themselves. Misconfiguration is one of the major contributors to issues related to cloud security, according to empirical evidence obtained through cloud security breaches conducted by the company IBM Security (Alzahrani et al., 2018; IBM Security, 2023).

However, this practice is unsustainable as the systems within the cloud environment continue to expand in size. Another potential security risk that can arise from a system hosted on a cloud environment is the existence of disparate security levels that can occur when the security department configures different systems within a specific product while ignoring the same configurations in another product within the same environment.

Rather than trying to circumvent strict restrictions, malicious actors prefer targeting inconsistencies in order to use them as opportunities. As much as the security of other platforms may be tight, one incorrect setup could be a gateway.

On the other hand, the configuration drift ought not be perceived as an operational bizarre rather, it functions as an integral result of fragmented governance. A security posture will

deteriorate silently unless constantly subject to validation, as such, multi-cloud environments will be in an insecure status perpetually.

4.0 Cross-Cloud Visibility Gaps

4.1 Fragmented Monitoring and Detection

Intensive details on system activities are required in designing an effective threat detection system. In multi-cloud computing, the applications of logging and monitoring have been identified to be scattered across different solutions provided by the service providers in different formats in varying levels of detail (Cloud Security Alliance, 2023). Security operations lack a clear understanding in this context of the system activities and potential threats.

Attackers spread their activity across these platforms to effectively exploit these blind spots. There can be signs of coordinated attack activity, presented as a group of harmless events, but these signs will not be recognized when central correlation is not performed. There is also the issue of differing levels of logging and alerting.

Security teams mostly rely on manual analysis or log aggregation that is gated by time, which extends the attacker dwell time and slows down detection (Mandiant, 2023; Verizon, 2024). This latency allows adversaries to create persistence, escalate privileges, and exfiltrate data well before containment.

To conclude, proactive defence is impaired in a disjointed surveillance network. Even a robust line of defence fails in a state of disarray on various cloud networks because detection occurs in a reactive state without a cohesive view of network operations.

4.2 Impact on Incident Response and Forensics

Further, forensic investigation, as well as event response, is hindered due to the presence of sight gaps in a multi-cloud environment. These procedures include the ability for the environment to detect an event, correlate the event, and log the event. These activities are undermined in a multi-cloud environment, particularly due to factors such as the presence of log formats, standardization of timestamps, as well as those relating to retention (NIST, 2018).

In case of an incident, it would be hard for a security team to replicate the attack timings across multiple cloud platforms, which would delay containment and create more uncertainty. Most importantly, fragmented visibility undermines post-event analysis, compromising organizational learning as well as the incidence of its recurrence.

Ineffective capabilities within the forensic environment, on the other hand, might lead to lack of regulatory compliance, as well as an absence of adequate, accountable support within the realm of governance, especially during a legal investigation or probe audit. In such a scenario, an organization might be unable to show diligence as a reliable source.

Significantly, the failure in handling incidences in hybrid environments is architectural rather than process-oriented. Without proper view, even well-trained teams are ineffective in building a structurally effective response.

In conclusion, issues with security in multi-cloud infrastructure are not coincidental but are deeply embedded in its architecture. Although identity-based attacks exploit federative trust models and privileged identity, policy inconsistency leads to vulnerabilities that are exploitable because of misconfigurations and changes. Limited visibility in cloud infrastructures makes response and detection less prompt. It has not resulted primarily from particular flaws in these infrastructures but from multi-cloud architectural complexity.

As per this particular research, irrespective of how sophisticated it may be in terms of mitigation capabilities, it cannot be assessed without identification with respect to what comprises any particular vector of threat. For instance, any issues identified with respect to what comprises a structure that houses multi-cloud are more likely to pose a threat to even the most advanced measures if identity governance, consistency of policy, and unification are not considered a part of the process. A strong foundation has thus emanated from this particular piece of research with regards to assessing sophisticated threats.

5.0 Zero Trust Architecture in Multi-Cloud Deployments

The Zero Trust architecture is extremely effective in multi-cloud deployments because it eliminates implicit trust and requires all users, devices, and workloads to be continuously checked before access is granted. In environments where resources are spread across several cloud service providers, Zero Trust decreases the risk of unwanted access by imposing tight identity verification and access control regulations independent of network location. Zero Trust restricts lateral movement within and across cloud platforms by micro-segmentation and granular policy enforcement, reducing the potential effect of a security compromise. Additionally, Zero Trust allows for consistent security policy enforcement and increased visibility across diverse contexts. However, its effectiveness is dependent on careful planning, integration with identity systems, and continual operational maturity, as bad implementation can exacerbate complexity and misconfiguration hazards.

6.0 Identity and Access Management (IAM)

Identity and Access Management (IAM) is critical for safeguarding multi-cloud settings because it centralizes management over user, service, and application identities across many clouds. Effective IAM implementations enforce least-privilege access via role-based or attribute-based access controls, dramatically lowering the possibility of credential abuse and privilege escalation. IAM also improves security by managing the entire identity lifecycle, which ensures that access rights are properly supplied, evaluated, and revoked as roles change. Furthermore, IAM's full logging and auditing features facilitate compliance requirements and incident investigations. Despite its importance, IAM effectiveness can be undermined by inconsistencies in identity federation, too permissive roles, and configuration problems that are typical in complicated multi-cloud environments.

7.0 Continuous Authentication

Continuous authentication improves threat prevention by constantly reviewing user and session behaviour throughout the access lifespan, as opposed to depending exclusively on initial login verification. This technique enables enterprises to dynamically analyse risk in multi-cloud

systems by leveraging contextual cues such as device posture, location, access patterns, and behavioural anomalies. Continuous authentication is especially useful for identifying compromised credentials, session hijacking, and insider threats because access can be restricted or terminated if suspicious behaviour is identified. By adhering to Zero Trust principles, it ensures continuous verification and adaptive access control. However, its successful implementation necessitates powerful analytics, high-quality telemetry, and careful tuning to avoid excessive false positives, which can harm user experience and productivity.

8.0 Automated Threat Detection

The scale, speed, and complexity of cloud workloads necessitate automated threat detection systems in multi-cloud installations. These systems collect and analyse security telemetry from numerous cloud platforms, resulting in centralized visibility and cross-environment threat correlation. Using machine learning and behavioural analytics, automated detection technologies can detect sophisticated threats, misconfigurations, and unusual activity that traditional signature-based solutions may miss. Automation greatly reduces the average time to detect threats, allowing for faster action in dynamic cloud environments. Nonetheless, the success of automated threat detection is strongly dependent on the quality of log data, proper integration with cloud-native services, and ongoing improvement of detection rules to decrease noise and false warnings.

9.0 Automated Response Mechanisms

Automated response systems enhance the effectiveness of threat mitigation in multi-cloud environments by allowing for speedy and consistent event containment actions. Automated solutions, using preset playbooks and orchestration procedures, can isolate compromised workloads, disable impacted accounts, and enforce network restrictions without requiring operator involvement. This decreases attacker dwell time and possible harm while also providing consistent incident response across several cloud platforms. Integration with SIEM and SOAR technologies improves coordination and scalability of response operations. However, excessive or poorly managed automation can pose operational risks such as service disruption or

inadvertent access revocation, highlighting the importance of human monitoring and staged reaction mechanisms.

10.0 Strengths of Advanced Threat Mitigation Approaches

The effectiveness of modern threat mitigation in multi-cloud environments is determined not by a single technology, but by an architectural transition from static perimeters to dynamic, data-driven decision-making.

10.1 Granularity of Identity-Centric Security

The most significant strength of modern advanced mitigation is the separation of security from network topology. Traditional security is based on "choke points" (firewalls), which are ineffective when resources are scattered across AWS, Azure, and on-premise data centers. Advanced identity-centric security approaches, particularly Workload Identity Federation (WIF), improve security by eliminating the requirement for long-lived credentials. Instead of static keys that can be stolen and used endlessly, new methods employ ephemeral, cryptographically verified tokens that expire automatically. This greatly limits the "blast radius" if a credential is compromised (Deochake et al., 2025). Unlike binary access (allow/deny), current IAM takes into account context, user location, device health, and behavior before giving access, making credential theft inadequate for a successful attack.

10.2 Predictive Capability of AI and ML

Artificial intelligence integration signifies a shift in the defense planning from reactive to proactive. Conventional signature-based detection requires a known "patient zero," which makes it ineffective against zero-day vulnerabilities. A "normality baseline" for user and API behavior is something that AI-driven platforms are excellent at. Systems can achieve detection rates surpassing 96% against complex Advanced Persistent Threats (APTs) by employing hybrid models, such as integrating XGBoost with Deep Neural Networks, to detect small deviations that imply low-and-slow data exfiltration or lateral movement (Imashev, 2025). AI systems are able to link seemingly unconnected events from many clouds, such as an Azure login failure and an AWS permission change, to identify them as a part of a single, well-planned assault campaign.

10.3 Operational Velocity via Automation (SOAR)

A crucial factor in multi-cloud security is speed. A multi-cloud stack generates so many alarms that human analysts are unable to physically examine them. Platforms for Security Orchestration, Automation, and Response (SOAR) convert manual playbooks into automated workflows. SOAR can reduce an attacker's "dwell time" from days to seconds by carrying out pre-authorized measures like isolating a virtual machine (VM) or revoking a user token in milliseconds when a high-fidelity threat (like ransomware behavior) is detected (Hassan et al., 2025).

11.0 Limitations and Critical Challenges

While modern technologies have strong capabilities, they create new levels of complexity that, if not managed properly, might paradoxically impair the security posture.

11.1 The "Black Box" Problem and Explainability

The inability of AI-driven threat intelligence to explain itself is a significant constraint. Deep Learning models frequently serve as "black boxes," which means they provide a likelihood of a hazard without explaining why. If an AI model flags a legal business process as malicious (false positive), and the security team does not understand the logic behind the flag, they may hesitate to prevent it, or block key business processes. This causes "alert fatigue," in which analysts start to overlook warnings (Imashev, 2025). Attackers are now adopting "model poisoning" approaches, which involve injecting bad data into the system in order to train the AI to disregard harmful activity, effectively blinding the defense.

11.2 Integration Complexity and Configuration Drift

The "Best-of-Breed" method of purchasing the best individual technology for network security, identity and access management, and endpoint protection results in a fragmented ecosystem. A threat discovered by a Google Cloud endpoint tool may not be relayed to the AWS firewall. This lack of compatibility leads to gaps in visibility. Thangaraju (2024) points out that maintaining consistent rules enforcement across diverse environments is impossible to do

manually. A policy update in one cloud may not replicate in another, leaving "orphaned" resources that are unprotected and unmonitored.

11.3 The Security-Performance Trade-off

The implementation of continuous, real-time inspection of all traffic and API requests creates latency. The overhead of decrypting communication, analyzing it using ML models, and re-encrypting it can significantly decrease performance in high-performance applications or resource-constrained IoT contexts (Alobaywi et al., 2025).

12.0 Future Directions and Emerging Trends

The future of multi-cloud security is evolving toward unified, self-regulating systems that require less human intervention and provide greater transparency.

12.1 Cybersecurity Mesh Architecture (CSMA)

In order to solve this problem of tool fragmentation, the industry is working toward Cybersecurity Mesh Architecture (CSMA). CSMA is a modular method that allows different security tools to be integrated into a single framework. This allows tools to dynamically share intelligence and coordinate reactions, regardless of which vendor built them. This results in a collaborative environment in which identity, context, and policy enforcement are separate from the underlying infrastructure, allowing for consistent security policies across all clouds and on-premise systems (Chandramouli & Butcher, 2023).

12.2 Autonomous Self-Healing Systems

Future SOAR iterations will evolve into Autonomous Security Systems. Instead of simply banning an IP address, these technologies will be able to "heal" the environment. If a vulnerability is discovered in a container, the system may automatically patch the image, rebuild it, and reinstall it without the need for human participation. This changes the goalposts from "Threat Response" to "Continuous Resilience."

12.3 Privacy-Preserving AI (Federated Learning)

As data protection restrictions (such as GDPR) tighten, transmitting all security logs to a centralized cloud for AI analysis becomes risky. Federated Learning is an emerging approach in which the AI model is sent to the data (at the edge), rather than the other way around. This enables enterprises to train advanced threat detection algorithms without ever revealing sensitive raw data or user records.

12.4 Quantum-Resistant Cryptography

With the arrival of quantum computing, conventional encryption protocols (RSA) will become outdated. To future-proof multi-cloud security, integrate Post-Quantum Cryptography (PQC) and blockchain-based immutable ledgers to ensure that identity verification and data integrity are secure even against quantum-level decoding attacks.

13.0 Conclusion

The migration to multi-cloud infrastructure marks a significant paradigm shift in enterprise technology, driven by the need for operational stability, scalability, and the avoidance of vendor lock-in. However, as this study has shown, the very variability that brings these economic benefits poses significant security issues. Traditional network perimeters have collapsed, exposing businesses to a complex threat landscape dominated by identity-based attacks, policy inconsistencies, and visibility holes. In a dispersed environment, the "perimeter" is clearly defined by Identity and Access Management (IAM) rather than a physical barrier.

The research study of advanced mitigation solutions, notably Zero Trust Architecture, Continuous Authentication, and AI-driven threat intelligence, found that static defenses are no longer sufficient. While these technologies provide strong capabilities for microsegmentation and proactive threat detection, their usefulness is usually hampered by operational complexity and tool fragmentation. The "siloed" structure of security measures among cloud providers creates significant blind spots, where a lack of interoperability allows intruders to remain unnoticed. As

a result, relying on "best-of-breed" point solutions is unsustainable in the absence of a unified governance structure to connect them together.

Finally, the future of multi-cloud security is determined not by the acceptance of a single superior technology, but by the transition to architectural unity. The rising themes of Cybersecurity Mesh Architecture (CSMA) and self-healing systems promise a future in which security is adaptable, modular, and mostly automated. To be resilient, companies must shift away from reactive posture management and toward proactive, identity-centric defense-in-depth methods. Success in this new era will necessitate a balance in which security measures are as dynamic and distributed as the cloud environments they are tasked with safeguarding, ensuring that governance evolves in tandem with innovation.

References

Akuthota, A. K. (2025). Securing Multi-Cloud Environments with SAP BTP, AI, and IAM: A Comprehensive Approach. *International Journal on Science and Technology*, 16(2), 1-6.

Alobaywi, B., Almutairi, M. G., & Sheldon, F. T. (2025). Multi-Tenancy Security in IoT-Cloud Systems: Challenges, Mitigations, and Future Directions. Preprints. <https://doi.org/10.20944/preprints202512.0463.v1>

Chandramouli, R., & Butcher, Z. (2023). A Zero Trust Architecture Model for Access Control in Cloud-Native Applications in Multi-Location Environments (NIST SP 800-207A). National Institute of Standards and Technology. <https://doi.org/10.6028/NIST.SP.800-207A>

Deochake, S., Murphy, R., & Gearheart, J. (2025). A Multi-Cloud Framework for Zero-Trust Workload Authentication. arXiv. <https://arxiv.org/abs/2510.16067>

Hassan, A., Rauf, A., Shafqat, N., Latif, R., & Khan, H. (2025). ZenGuard a machine learning based zero trust framework for context aware threat mitigation using SIEM SOAR and UEBA. *Scientific Reports*, 15(35871). <https://doi.org/10.1038/s41598-025-20998-4>

Imashev, A. (2025). AI-Driven Zero-Trust Security Framework for Detecting Advanced Persistent Threats in Cloud Environments. *IRE Journals*, 9(5), 477-485.

Thangaraju, V. (2024). Security Considerations in Multi-Cloud Environments with Seamless Integration: A Review of Best Practices and Emerging Threats. *Transactions on Engineering and Computing Sciences*, 12(02), 226-238.

Akuthota, A. K. (2025). Securing Multi-Cloud Environments with SAP BTP, AI, and IAM: A Comprehensive Approach. *International Journal on Science and Technology*, 16(2), 1–6.

National Institute of Standards and Technology. (2020). Zero trust architecture (NIST Special Publication 800-207). *U.S. Department of Commerce*, 1–56.

National Institute of Standards and Technology. (2017). Digital identity guidelines (NIST Special Publication 800-63). *U.S. Department of Commerce*, 1–128.

Cloud Security Alliance. (2021). Security guidance for critical areas of focus in cloud computing (Version 4.0). *Cloud Security Alliance*, 1–175.

IBM Security. (2022). Cost of a data breach report. *IBM Corporation*, 1–67.

European Union Agency for Cybersecurity. (2021). ENISA threat landscape for cloud computing. *ENISA Reports*, 1–84.

Alzahrani, N., Bulusu, N., & Thampi, S. M. (2018). Cloud security: A survey of threats and solutions. *Journal of Network and Computer Applications*, 104, 1–15.
<https://doi.org/10.1016/j.jnca.2017.12.004>

Behl, A., & Behl, K. (2017). *Cyberwar: The next threat to national security and what to do about it*. Oxford University Press.

Cloud Security Alliance. (2023). *Security guidance for critical areas of focus in cloud computing v4.0*. <https://cloudsecurityalliance.org>

CISA. (2023). *Securing cloud business applications*. <https://www.cisa.gov>

ENISA. (2021). *Cloud security for SMEs: Challenges and recommendations*. <https://www.enisa.europa.eu>

IBM Security. (2023). *Cost of a data breach report 2023*. <https://www.ibm.com/security>

ISO/IEC. (2022). *ISO/IEC 27001:2022 Information security management systems*. International Organization for Standardization.

Mandiant. (2023). *M-Trends 2023: Insights into today's breach landscape*. <https://www.mandiant.com>

NIST. (2018). *Guide to integrating forensic techniques into incident response (SP 800-86)*. National Institute of Standards and Technology.

NIST. (2023). *Security and privacy controls for information systems (SP 800-53 Rev. 5)*. National Institute of Standards and Technology.

Rose, S., Borchert, O., Mitchell, S., & Connelly, S. (2020). *Zero Trust Architecture*. NIST Special Publication 800-207.

Verizon. (2024). *Data breach investigations report*. <https://www.verizon.com/business/resources>

AWS. (2023). *AWS security best practices for multi-account environments*. <https://docs.aws.amazon.com>

Microsoft. (2023). *Azure security benchmark*. <https://learn.microsoft.com>

Google Cloud. (2023). *BeyondCorp enterprise security model*. <https://cloud.google.com>

Hashizume, K., Rosado, D. G., Fernández-Medina, E., & Fernandez, E. B. (2013). An analysis of security issues for cloud computing. *Journal of Internet Services and Applications*, 4(1), 1–13. <https://doi.org/10.1186/1869-0238-4-5>

Zhang, Q., Chen, M., Li, L., & Li, Y. (2018). Security challenges in multi-cloud computing. *IEEE Access*, 6, 40608–40620. <https://doi.org/10.1109/ACCESS.2018.2852618>

Pearson, S. (2013). Privacy, security and trust in cloud computing. *Privacy and Security for Cloud Computing*, 3–42. Springer.

Gartner. (2023). *Market guide for cloud-native application protection platforms*.